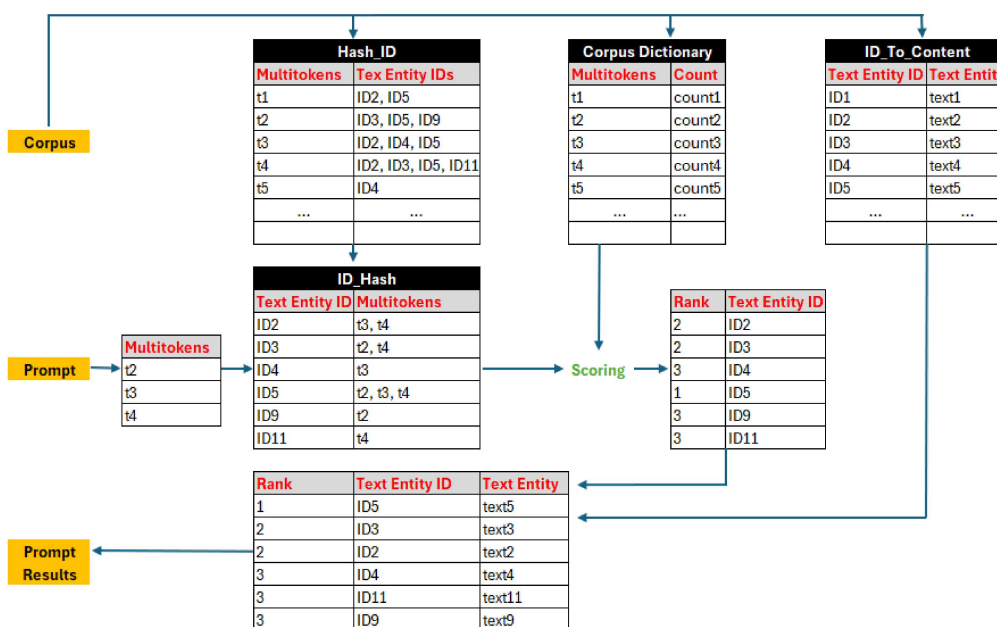
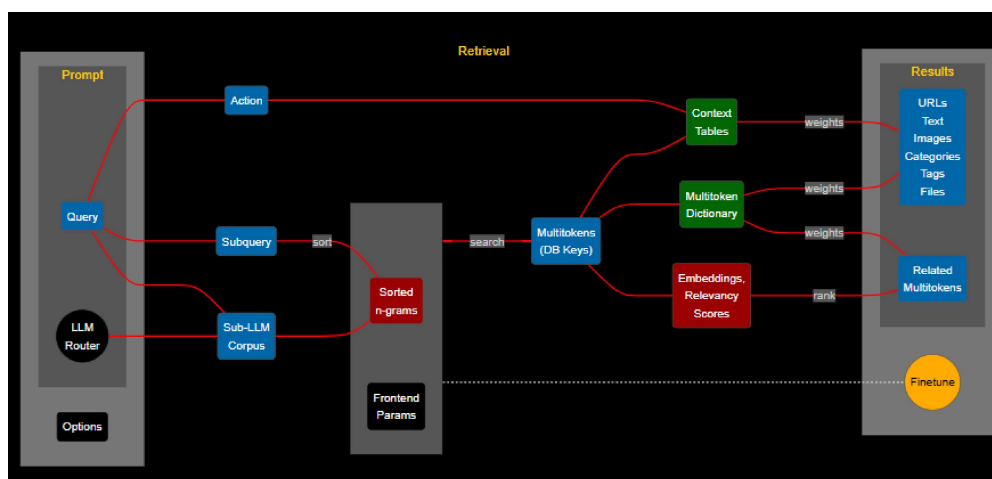


Building Disruptive AI & LLM Technology From Scratch



Contents

I	Hallucination-Free LLM with Real-Time Fine-Tuning	6
1	Enterprise xLLM for Search & Retrieval	7
1.1	Efficient database architecture: nested hashes	7
1.2	From frontend prompts to backend tables	8
1.3	What is not covered here	9
2	Parameters, features, and fine-tuning	11
2.1	Backend parameters	11
2.2	Frontend parameters	12
2.3	Agents	12
2.4	Reproducibility	13
2.5	Singularization, stemming, auto-correct	14
2.6	Augmentation, distillation, and frontend tables	14
2.7	In-memory database, latency, and scalability	15
3	Case study	17
3.1	Real-time fine-tuning, prompts, and command menu	18
3.2	Sample session	19
3.3	Web API for enterprise xLLM	29
3.3.1	Left panel: command menu and prompt box	29
3.3.2	Right panel: prompt results	29
3.3.3	Next steps	31
3.4	Conclusions and references	31
4	Appendix	32
4.1	Python code	32
4.2	Thirty features to boost LLM performance	44
4.2.1	Fast search and caching	44
4.2.2	Leveraging sparse databases	45
4.2.3	Contextual tokens	45
4.2.4	Adaptive loss function	45
4.2.5	Contextual tables	45
4.2.6	Smart crawling	45
4.2.7	LLM router, sub-LLMs, and distributed architecture	45
4.2.8	From one trillion parameters down to two	46
4.2.9	Agentic LLMs	46
4.2.10	Data augmentation via dictionaries	46
4.2.11	Distillation done smartly	46
4.2.12	Reproducibility	47
4.2.13	Explainable AI	47
4.2.14	No training, in-memory LLM	47
4.2.15	No neural network	47
4.2.16	Show URLs and references	48
4.2.17	Taxonomy-based evaluation	48
4.2.18	Augmentation via prompt data	48
4.2.19	Variable-length embeddings, indexing, and database optimization	48
4.2.20	Favor backend over frontend engineering	48
4.2.21	Use NLP and Python libraries with caution	49
4.2.22	Self-tuning and customization	49

4.2.23	Local, global parameters, and debugging	49
4.2.24	Displaying relevancy scores, and customizing scores	50
4.2.25	Intuitive hyperparameters	50
4.2.26	Sorted n -grams and token order preservation	50
4.2.27	Blending standard tokens with tokens from the knowledge graph	50
4.2.28	Boosted weights for knowledge-graph tokens	51
4.2.29	Versatile command prompt	51
4.2.30	Boost long multitokens and rare single tokens	51
4.2.31	Disambiguation	51
4.3	LLM glossary	51
II	Outperforming Neural Nets and Classic AI	55
5	Building and evaluating a taxonomy-enriched LLM	57
5.1	Project and solution	59
5.2	Python code	60
6	LLM for Clustering and Predictions	66
6.1	Project and solution	67
6.2	Visualizations and discussion	69
6.3	Python code	70
7	Fast, High-Quality NoGAN for Data Synthetization	78
7.1	Project description	78
7.2	Solution	80
7.3	Python implementation	81
8	Hierarchical Bayesian NoGAN for Data Synthetization	87
8.1	Methodology	87
8.1.1	Base algorithm	88
8.1.2	Loss function	88
8.1.3	Hyperparemeters and convergence	89
8.1.4	Acknowledgments	90
8.2	Case studies	90
8.2.1	Synthesizing the student dataset	91
8.2.2	Synthesizing the Telecom dataset	93
8.2.3	Other case studies	94
8.2.4	Auto-tuning the hyperparameters	95
8.2.5	Evaluation with multivariate ECDF and KS distance	96
8.3	Conclusion	97
8.4	Python implementation	98
9	Boosting Model Evaluation with Smart Adaptive Loss Functions	107
9.1	Project and solution	108
9.2	Python code	112
III	Innovations in Statistical AI	117
10	Building a Ranking System to Enhance Prompt Results	119
10.1	Relevancy scores and rankings	120
10.2	Case study	121
10.2.1	xLLM for auto-indexing, cataloging and glossary generation	121
10.2.2	xLLM for scientific research	122
10.3	Python code	123
10.4	Appendix: smart ranking in a nutshell	129
11	Probabilistic Search: Alternative to Vector Search	131
11.1	Motivation and architecture	132
11.2	Applications	133
11.2.1	Embeddings and large language models	133

11.2.2	Generating and evaluating synthetic data	134
11.2.3	Clustering, dataset comparisons, outlier detection	135
11.3	Project and solution	136
11.4	Python code	137
12	Strong Random Generators for Reproducible AI	141
12.1	Strong Randomness and reproducibility: two key components	141
12.2	Computing the digits of special math constants	142
12.2.1	P-adic valuations	142
12.2.2	Digit blocks, speed of convergence	143
12.2.3	A plethora of interesting pseudo-random sequences	143
12.3	Testing random number generators	144
12.3.1	Theoretical properties of the digits of $\sqrt{2}$	144
12.3.2	Fast recursion and congruential equidistribution	145
12.3.3	Exponential system: predicting the next block	147
12.4	Python code	149
12.4.1	Fast recursion	149
12.4.2	Main code	150
13	Sampling Outside the Observation Range	155
13.1	Quantile convolution	155
13.2	Truncated Gaussian mixtures and bias detection	157
13.3	Case studies	157
13.3.1	Conclusion	158
13.4	Python code	159
14	Additional Resources	162
14.1	Accelerating convergence of parameter estimates	162
14.1.1	First case study	162
14.1.2	Second case study	162
14.2	Generic, all-in-one curve fitting, regression and clustering	164
14.2.1	Solution, R-squared and backward compatibility	165
14.2.2	Model upgrades	166
14.2.3	Case studies	166
14.2.4	Python code	170
14.3	A simple geospatial interpolation method	171
14.3.1	Problem in two dimensions	171
14.3.2	Spatial interpolation of the temperature dataset	173
14.3.3	Python code	174
14.4	Math-free, parameter-free gradient descent	174
14.4.1	Introduction	175
14.4.2	Implementation details	175
14.4.3	General comments about the methodology and parameters	178
14.4.4	Mathematical version of gradient descent and orthogonal trajectories	179
14.4.5	Python code	180
15	Trading the S&P 500 Index	181
15.1	Non-standard strategies	182
15.2	Selecting a strategy based on its features	183
15.3	Python code and dataset: 40 years' worth of historical data	184
	Bibliography	188
	Index	190

Introduction

This book features new advances in game-changing AI and LLM technologies built by GenAITECHLab.com. Written in simple English, it is best suited for engineers, developers, data scientists, analysts, consultants and anyone with an analytic background interested in starting a career in AI. The emphasis is on scalable enterprise solutions, easy to implement, yet outperforming vendors both in terms of speed and quality, by several orders of magnitude.

Each topic comes with GitHub links, full Python code, datasets, illustrations, and real life case studies, including from Fortune 100 company. Some of the material is presented as enterprise projects with solution, to help you build robust applications and boost your career. You don't need expansive GPU and cloud bandwidth to implement them: a standard laptop works.

Part I focuses on high performance in-memory agentic multi-LLMs for professional users and enterprise, with real-time fine-tuning, self-tuning, no weight, no training, no latency, no hallucinations, no GPU. Made from scratch, leading to replicable results, leveraging explainable AI, adopted by Fortune 100. With a focus on delivering concise, exhaustive, relevant, and in-depth search results, references, and links. See also the section on 31 features to substantially boost RAG/LLM performance.

Part II discusses related large-scale systems also benefiting from a light-weight but more efficient architecture. It features LLMs for clustering, classification, and taxonomy creation, leveraging the knowledge graphs embedded in and retrieved from the input corpus when crawling. Then, in chapters 7 and 8, I focus on tabular data synthetization, presenting techniques such as NoGAN, that significantly outperform neural networks, along with the best evaluation metric. The methodology in chapter 9 applies to most AI problems. It offers a generic tool to improve any existing architecture relying on gradient descent, such as deep neural networks.

Part III features a collection of methods that you can integrate in any AI system to boost performance. Based on a modern approach to statistical AI, they cover probabilistic vector search, sampling outside the observation range, strong random number generators, math-free gradient descent, beating the slow statistical convergence of parameter estimates dictated by the Central Limit Theorem, exact geospatial interpolation for non-smooth systems, and more. Efficient chunking and indexing for LLMs is the topic of chapter 10. Finally, chapter 15 shows how to optimize trading strategies to consistently outperform the stock market.

About the author

Vincent Granville is a pioneering GenAI scientist and machine learning expert, co-founder of Data Science Central (acquired by a publicly traded company in 2020), Chief AI Scientist at MLTechniques.com, former VC-funded executive, author and patent owner – one related to LLM. Vincent's past corporate experience includes Visa, Wells Fargo, eBay, NBC, Microsoft, and CNET.



Vincent is also a former post-doc at Cambridge University, and the National Institute of Statistical Sciences (NISS). He published in *Journal of Number Theory*, *Journal of the Royal Statistical Society* (Series B), and *IEEE Transactions on Pattern Analysis and Machine Intelligence*. He is the author of multiple books, available [here](#), including “Synthetic Data and Generative AI” (Elsevier, 2024). Vincent lives in Washington state, and enjoys doing research on stochastic processes, dynamical systems, experimental math and probabilistic number theory.

Index

- χ^2 distribution, 147
- k -NN, 52, 132
- k -means clustering, 67
- k -medoids clustering, 67
- n -gram, 8
 - sorted, 50
- p -adic valuation, 142
- p -value, 147, 149

- abbreviation dictionary, 14
- acronyms, 46, 52
- action, 8, 12, 51
- activation function, 97
- Adam (stochastic gradient descent), 52
- agent, 12, 18, 29, 46, 53, 119, 130
 - multi-agent, 51
- ANN (approximate nearest neighbors search), 44, 54, 134
 - probabilistic ANN (pANN), 52, 134
- augmentation, 15, 48, 52
 - augmented knowledge graph, 15
 - taxonomy augmentation, 15, 122
- auto-correct, 14
- auto-encoder, 52
- auto-regressive model, 53
- auto-tuning, 89, 95, 97

- backend, 29, 52, 53, 119
- base, 146
- basin of attraction, 182
- batch, 89, 91, 97
- Bayesian hierarchical models, 90
- benchmarking, 52
- Beurling primes, 164
- binary search
 - interpolated, 135
- binning, 108
- bit shift, 145
- bootstrapping, 69, 166

- caching, 44, 54
- cataloging, 122
- categorical feature, 89
- CDF regression, 167
- central limit theorem, 162
- chunking, 121, 129
- confidence interval, 69, 162
 - model-free, 81
- confidence region, 166
- congruential equidistribution
 - asymptotic, 145
- connected components, 67
- context, 119
 - contextual pairs, 45
 - contextual tables, 45
 - contextual tokens, 45
- contour level, 175, 179
- contour lines, 175
- convergence
 - absolute, 171
 - conditional, 171
- convolution product, 155
- copula, 90, 98
- cosine similarity, 54, 134
- covariance matrix, 110
- Cramér's V, 89
- crawling
 - smart crawling, 8, 45
- critical line (number theory), 175
- cross-validation, 48, 68, 87, 166
- customization (LLM), 49

- data distillation, 79
- database
 - graph database, 53
 - JSON, 48
 - key-value pairs, 48
 - nested hash, 48
 - vector database, 48
- debugging, 49
- deep neural network (DNN), 52, 97, 107
- dendrogram, 69
- diffusion, 52, 80
- diffusion model, 155
- digit block, 143
- digit-preserving, 143
- dimensionality reduction, 166
- Dirichlet character, 164
- Dirichlet- L function, 164
- disembiguation, 51
- distance matrix, 66
- distillation, 14, 46, 52
- distributed computing, 46
- distribution
 - logistic, 166
- DNN (deep neural network), 45
- dot product, 133, 165
- dummy variable, 89
- dynamical system, 182

- ECDF (empirical distribution), 54, 78, 88, 107
- embedding, 19, 45, 52

- quantized embedding, 48
- variable length embedding, 7, 48, 133
- empirical density function, 96
- empirical distribution, 67, 167
 - multivariate, 54, 78, 87, 90, 96, 107, 135
- encoding
 - smart encoding, 68
- EPDF (empirical probability density function), 107
- epoch (neural networks), 132, 137
- equidistribution modulo 1, 146
- Euler product, 164
- evaluation, 107
- evaluation metrics, 48, 52, 53
- exhaustivity (LLM/RAG output), 119
- explainable AI, 50, 52, 78, 87, 97
- feature encoding, 54
- feature engineering, 53
- feature selection, 166
- fine-tuning, 52, 130
 - exhaustive results, 47, 119
 - in real time, 17
 - LLM parameters, 46
 - self-tuning, 49
- frontend, 52
- GAN (generative adversarial network), 13, 78, 107, 155
 - hierarchical GAN, 90
 - NoGAN, 14
- Gaussian mixture model, 110, 155
- generative adversarial network, 52, 80, 87, 134, 155
- GIS, 173
- GPT, 52
- GPU, 47
- gradient descent, 69, 88, 97, 110, 132, 179, 181
 - steepest, 136
 - stochastic, 107, 136, 137
- gradient operator, 165
- graph
 - graph database, 53
 - knowledge graph, 48
- GRH (Generalized Riemann Hypothesis), 164
- grid search, 79, 181
 - smart grid search, 47, 111, 181
- Hadamard product, 88
- hallucination, 48
- hash table
 - inversion, 66
 - nested hash, 8, 45, 46, 48, 53, 66, 119, 122
- Hellinger distance, 96, 107
 - multivariate, 45
- hierarchical Bayesian model, 98
- hierarchical clustering, 67
- hierarchical deep resampling, 87, 90
- holdout, 90
- Hungarian algorithm, 107, 132
- hyperparameter, 47, 50, 52, 79, 89
- hyperrectangles, 78
- identifiability (statistics), 131
- imputation (missing values), 90
- in-memory database, 16
- in-memory LLM, 16, 47
- indexing, 129
 - auto-indexing, 122
 - glossary generation, 122
 - text entity index, 48, 119
- integer division, 145
- integral probability metrics, 96
- interpolation, 182
- JSON, 8, 45
 - database, 48, 53
- kernel, 155
- kernel density estimation, 155
- key-value database, 8, 53
- knowledge graph, 48, 119, 129
- Kolmogorov-Smirnov distance, 67, 78, 89, 90, 96, 109, 135
 - multivariate, 45, 54, 107
- kriging, 171
- label, 95
- Lagrange multiplier, 165
- LangChain, 53
- large language model (LLM)
 - evaluation, 122
- large language models (LLM), 57
- Lasso regression, 165
- latency, 15
- law of the iterated logarithm, 147
- learning rate, 177
- link function, 167
- LLaMA, 45, 53
- LLM (large language model)
 - agentic LLM, 12
 - debugging, 49
 - for auto-indexing, 122
 - for cataloging, 122
 - for clustering, 13
 - for glossary generation, 122
 - for predictive analytics, 13
 - in-memory, 16
 - LLM router, 8, 46, 119
 - multi-LLM, 17
 - self-tuned, 49
 - sub-LLM, 17, 119
 - xLLM, 17
- logistic distribution, 166
- logistic regression, 166
- loss function, 67, 69, 88, 89, 97, 107, 132, 136, 181
 - adaptive loss, 45, 107
- matrix
 - positive semidefinite, 110
- mean squared error, 165
- Mersenne twister, 141, 144
- mixture model, 96
- mixture of experts, 46, 49, 119
- model identifiability, 165
- modulus (complex number), 175
- multi-agent system, 51, 53

- multimodal system, 53
- multinomial distribution, 78
- multiple root, 171
- multitoken (see token), 52
- nearest neighbors
 - K -NN, 131
 - approximate (ANN), 131
 - probabilistic (pANN), 131
 - search, 132
- NLG (natural language generation), 53
- NLP (natural language processing), 14, 53
- node (interpolation), 96
- NoGAN, 52, 78, 108, 134
 - constrained, 108
 - probabilistic, 108
- normal number, 146
- normalization, 53, 54
- notebookML, 122
- OpenAI, 110
- orthogonal trajectory, 175
- overfitting, 47, 67, 78, 165, 181
- parameter
 - in neural networks, 53, 54
 - in xLLM, 53
- parametric bootstrap, 170
- partial derivative, 171
- period, 142
- PMI (pointwise mutual information), 7, 9, 19, 31, 49, 50
- pointwise mutual information (PMI), 54, 133
- Poisson process, 137
- positive semidefinite (matrix), 110
- prediction interval, 166
- prefix, 142
- principal component analysis, 90
- PRNG (pseudo-random number generator), 14, 47, 141
- probability density function, 155
- prompt
 - command prompt options, 48
 - prompt compression, 134
 - prompt engineering, 48
 - prompt synthesizer, 122
- Python library
 - GenAI-evaluation, 109
 - Sklearn, 67
 - Sklearn-extra, 67
- quantile, 79, 109
 - extrapolated, 157
- quantile function, 91, 155
- quantile regression, 166
- quantization, 48
- R-squared, 165, 166
- Rademacher function, 164
- radix search, 135
- RAG (retrieval augmentation generation), 15, 53, 119
- random walk
 - simple, 147
- rank
 - relevancy rank, 120
- regularization, 54
- reinforcement learning, 49, 54, 97
- rejection sampling, 157
- relevancy
 - relevancy score, 47, 50, 119, 130
- replicability, 47
- reproducibility, 14, 47
- resampling, 108
- residue, 145
- retrieval, 48
 - retrieval augmentation generation (RAG), 53
- Riemann hypothesis
 - generalized, 164
- Riemann zeta function, 164, 175
- run test, 145
- saddle point, 177
- scalability, 15
- score
 - relevancy rank, 120
 - relevancy score, 29, 47
 - smart scoring, 130
- search
 - probabilistic search, 48
 - vector search, 48
- seed (random number generator), 13, 89
- self-tuning, 52, 54
- separator (text), 54
- similarity metric, 67
- simulated annealing, 97
- singularization, 14
- steepest descent, 179
- stemming, 14
- stopwords, 8, 49
- swap, 88, 97
- synonyms dictionary, 14
- synthetic data, 54, 107, 157, 171, 174
 - constrained, 108, 111
- synthetic function, 164
- taxicab distance, 137
- taxonomy, 53
 - augmentation, 122
 - taxonomy creation, 57, 122
- tensor, 89
- TensorFlow, 89
- text entity, 12, 29, 46, 48, 52, 119
 - sub-entity, 18
 - text entity ID, 29
- time complexity, 79
- token, 52, 54
 - contextual, 45, 66
 - graph token, 119, 121, 130
 - multitoken, 7, 19, 52, 66, 119
- training, 47, 52
- transformer, 13, 45, 47, 52, 54, 97
- truncated Gaussian, 157
- unit ball, 137

- validation set, [78](#), [90](#), [166](#)
- vanishing gradient, [89](#)
- variational auto-encoder, [52](#)
- vector database, [48](#), [54](#)
- vector search, [48](#), [54](#), [131](#)
- vectorization, [79](#)

- Wasserstein GAN (WGAN), [97](#), [107](#)
- Weibull distribution, [137](#)
- weighted regression, [166](#)
- Wiener process, [147](#)

- XGboost, [78](#)
- xLLM, [17](#), [57](#)
 - e-xLLM (for enterprise corpus), [119](#)