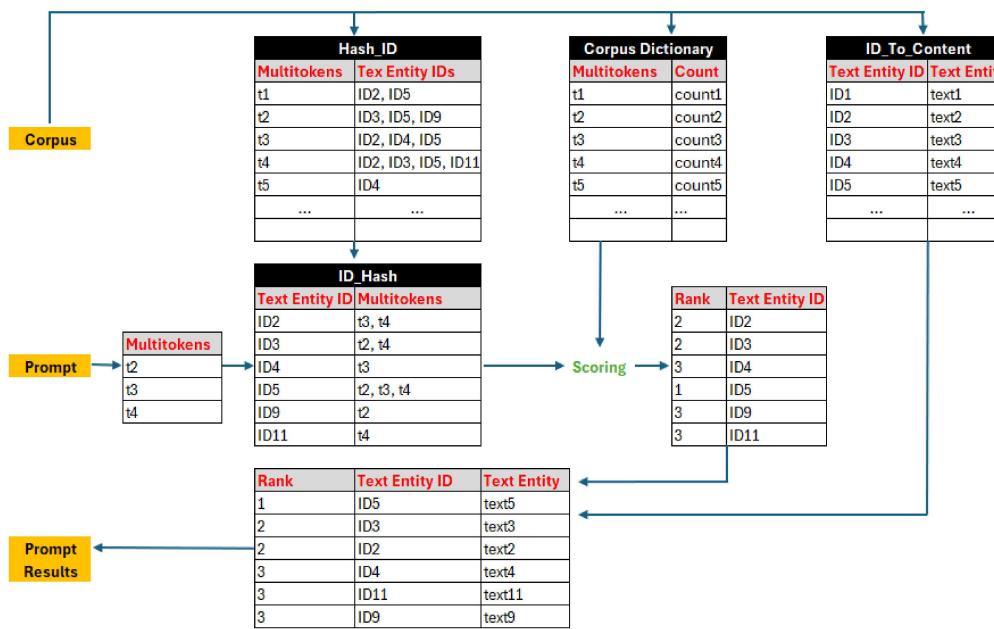
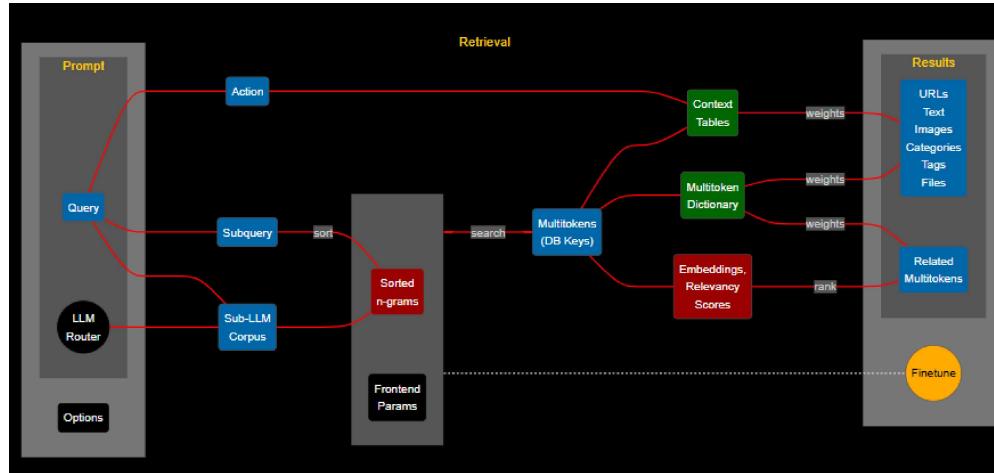


Building Disruptive AI & LLM Technology From Scratch



Contents

I Hallucination-Free LLM with Real-Time Fine-Tuning	6
1 Enterprise xLLM for Search & Retrieval	7
1.1 Efficient database architecture: nested hashes	7
1.2 From frontend prompts to backend tables	8
1.3 What is not covered here	9
2 Parameters, features, and fine-tuning	11
2.1 Backend parameters	11
2.2 Frontend parameters	12
2.3 Agents	12
2.4 Reproducibility	13
2.5 Singularization, stemming, auto-correct	14
2.6 Augmentation, distillation, and frontend tables	14
2.7 In-memory database, latency, and scalability	15
3 Case study	17
3.1 Real-time fine-tuning, prompts, and command menu	18
3.2 Sample session	19
3.3 Web API for enterprise xLLM	29
3.3.1 Left panel: command menu and prompt box	29
3.3.2 Right panel: prompt results	29
3.3.3 Next steps	31
3.4 Conclusions and references	31
4 Appendix	32
4.1 Python code	32
4.2 Thirty features to boost LLM performance	44
4.2.1 Fast search and caching	44
4.2.2 Leveraging sparse databases	45
4.2.3 Contextual tokens	45
4.2.4 Adaptive loss function	45
4.2.5 Contextual tables	45
4.2.6 Smart crawling	45
4.2.7 LLM router, sub-LLMs, and distributed architecture	45
4.2.8 From one trillion parameters down to two	46
4.2.9 Agentic LLMs	46
4.2.10 Data augmentation via dictionaries	46
4.2.11 Distillation done smartly	46
4.2.12 Reproducibility	47
4.2.13 Explainable AI	47
4.2.14 No training, in-memory LLM	47
4.2.15 No neural network	47
4.2.16 Show URLs and references	48
4.2.17 Taxonomy-based evaluation	48
4.2.18 Augmentation via prompt data	48
4.2.19 Variable-length embeddings, indexing, and database optimization	48
4.2.20 Favor backend over frontend engineering	48
4.2.21 Use NLP and Python libraries with caution	49
4.2.22 Self-tuning and customization	49

4.2.23 Local, global parameters, and debugging	49
4.2.24 Displaying relevancy scores, and customizing scores	50
4.2.25 Intuitive hyperparameters	50
4.2.26 Sorted n -grams and token order preservation	50
4.2.27 Blending standard tokens with tokens from the knowledge graph	50
4.2.28 Boosted weights for knowledge-graph tokens	51
4.2.29 Versatile command prompt	51
4.2.30 Boost long multitokens and rare single tokens	51
4.2.31 Disambiguation	51
4.3 LLM glossary	51
II Outperforming Neural Nets and Classic AI	55
5 Building and evaluating a taxonomy-enriched LLM	57
5.1 Project and solution	59
5.2 Python code	60
6 LLM for Clustering and Predictions	66
6.1 Project and solution	67
6.2 Visualizations and discussion	69
6.3 Python code	70
7 Fast, High-Quality NoGAN for Data Synthetization	78
7.1 Project description	78
7.2 Solution	80
7.3 Python implementation	81
8 Hierarchical Bayesian NoGAN for Data Synthetization	87
8.1 Methodology	87
8.1.1 Base algorithm	88
8.1.2 Loss function	88
8.1.3 Hyperparameters and convergence	89
8.1.4 Acknowledgments	90
8.2 Case studies	90
8.2.1 Synthesizing the student dataset	91
8.2.2 Synthesizing the Telecom dataset	93
8.2.3 Other case studies	94
8.2.4 Auto-tuning the hyperparameters	95
8.2.5 Evaluation with multivariate ECDF and KS distance	96
8.3 Conclusion	97
8.4 Python implementation	98
9 Boosting Model Evaluation with Smart Adaptive Loss Functions	107
9.1 Project and solution	108
9.2 Python code	112
III Innovations in Statistical AI	117
10 Building a Ranking System to Enhance Prompt Results	119
10.1 Relevancy scores and rankings	120
10.2 Case study	121
10.2.1 xLLM for auto-indexing, cataloging and glossary generation	121
10.2.2 xLLM for scientific research	122
10.3 Python code	123
10.4 Appendix: smart ranking in a nutshell	129
11 Probabilistic Search: Alternative to Vector Search	131
11.1 Motivation and architecture	132
11.2 Applications	133
11.2.1 Embeddings and large language models	133

11.2.2 Generating and evaluating synthetic data	134
11.2.3 Clustering, dataset comparisons, outlier detection	135
11.3 Project and solution	136
11.4 Python code	137
12 Strong Random Generators for Reproducible AI	141
12.1 Strong Randomness and reproducibility: two key components	141
12.2 Computing the digits of special math constants	142
12.2.1 P-adic valuations	142
12.2.2 Digit blocks, speed of convergence	143
12.2.3 A plethora of interesting pseudo-random sequences	143
12.3 Testing random number generators	144
12.3.1 Theoretical properties of the digits of $\sqrt{2}$	144
12.3.2 Fast recursion and congruential equidistribution	145
12.3.3 Exponential system: predicting the next block	147
12.4 Python code	149
12.4.1 Fast recursion	149
12.4.2 Main code	150
13 Sampling Outside the Observation Range	155
13.1 Quantile convolution	155
13.2 Truncated Gaussian mixtures and bias detection	157
13.3 Case studies	157
13.3.1 Conclusion	158
13.4 Python code	159
14 Additional Resources	162
14.1 Accelerating convergence of parameter estimates	162
14.1.1 First case study	162
14.1.2 Second case study	162
14.2 Generic, all-in-one curve fitting, regression and clustering	164
14.2.1 Solution, R-squared and backward compatibility	165
14.2.2 Model upgrades	166
14.2.3 Case studies	166
14.2.4 Python code	170
14.3 A simple geospatial interpolation method	171
14.3.1 Problem in two dimensions	171
14.3.2 Spatial interpolation of the temperature dataset	173
14.3.3 Python code	174
14.4 Math-free, parameter-free gradient descent	174
14.4.1 Introduction	175
14.4.2 Implementation details	175
14.4.3 General comments about the methodology and parameters	178
14.4.4 Mathematical version of gradient descent and orthogonal trajectories	179
14.4.5 Python code	180
15 Trading the S&P 500 Index	181
15.1 Non-standard strategies	182
15.2 Selecting a strategy based on its features	183
15.3 Python code and dataset: 40 years' worth of historical data	184
Bibliography	188
Index	190

Introduction

This book features new advances in game-changing AI and LLM technologies built by [GenAItchLab.com](#). Written in simple English, it is best suited for engineers, developers, data scientists, analysts, consultants and anyone with an analytic background interested in starting a career in AI. The emphasis is on scalable enterprise solutions, easy to implement, yet outperforming vendors both in terms of speed and quality, by several orders of magnitude.

Each topic comes with GitHub links, full Python code, datasets, illustrations, and real life case studies, including from Fortune 100 company. Some of the material is presented as enterprise projects with solution, to help you build robust applications and boost your career. You don't need expensive GPU and cloud bandwidth to implement them: a standard laptop works.

Part I: Hallucination-Free LLM with Real-Time Fine-Tuning focuses on high performance in-memory agentic multi-LLMs for professional users and enterprise, with real-time fine-tuning, self-tuning, no weight, no training, no latency, no hallucinations, no GPU. Made from scratch, leading to replicable results, leveraging explainable AI, adopted by Fortune 100. With a focus on delivering concise, exhaustive, relevant, and in-depth search results, references, and links. See also the section on 31 features to substantially boost RAG/LLM performance.

Part II: Outperforming Neural Nets and Classic AI discusses related large-scale systems also benefiting from a light-weight but more efficient architecture. It features LLMs for clustering, classification, and taxonomy creation, leveraging the knowledge graphs embedded in and retrieved from the input corpus when crawling. Then, in chapters 7 and 8, I focus on tabular data synthetization, presenting techniques such as No-GAN, that significantly outperform neural networks, along with the best evaluation metric. The methodology in chapter 9 applies to most AI problems. It offers a generic tool to improve any existing architecture relying on gradient descent, such as deep neural networks.

Part III: Innovations in Statistical AI features a collection of methods that you can integrate in any AI system to boost performance. Based on a modern approach to statistical AI, they cover probabilistic vector search, sampling outside the observation range, strong random number generators, math-free gradient descent, beating the slow statistical convergence of parameter estimates dictated by the Central Limit Theorem, exact geospatial interpolation for non-smooth systems, and more. Efficient chunking and indexing for LLMs is the topic of chapter 10. Finally, chapter 15 shows how to optimize trading strategies to consistently outperform the stock market.

About the author

Vincent Granville is a pioneering GenAI scientist and machine learning expert, co-founder of Data Science Central (acquired by a publicly traded company in 2020), Chief AI Scientist at [MLTechniques.com](#), former VC-funded executive, author and patent owner – one related to LLM. Vincent's past corporate experience includes Visa, Wells Fargo, eBay, NBC, Microsoft, and CNET.



Vincent is also a former post-doc at Cambridge University, and the National Institute of Statistical Sciences (NISS). He published in *Journal of Number Theory*, *Journal of the Royal Statistical Society* (Series B), and *IEEE Transactions on Pattern Analysis and Machine Intelligence*. He is the author of multiple books, available [here](#), including "Synthetic Data and Generative AI" (Elsevier, 2024). Vincent lives in Washington state, and enjoys doing research on stochastic processes, dynamical systems, experimental math and probabilistic number theory.

Chapter 15

Trading the S&P 500 Index

Before telling what this chapter is about, I want to discuss Figure 15.1. It features problems found in almost all AI systems. The solutions are relevant to many contexts well beyond FinTech. Each dot in the scatterplots represents the performance of an algorithm based on two parameters: β (X-axis) and σ (Y-axis). There is an extra parameter ω , with two values tested here: $\omega = 40$ on the left, and $\omega = 60$ on the right. The framework is highly non-linear, and indeed quite chaotic.

The color indicates the performance level: blue for great, green for good, gray for similar to baseline, orange for poor, and red for bad. The goal is to find decent parameter vectors (β, σ, ω) in the parameter space, based on the 128 tested configurations: 64 on the left, and another 64 on the right. Each test consists of running an algorithm on 12 different large portions of the input dataset, itself consisting of 40 years' worth of daily stock market data. In short, each performance measurement (color) is averaged over 12 different input datasets. The approach is similar to using [grid search](#) to optimize hyperparameters in a neural network. But there is more to it, applicable to the core of almost all neural networks and machine learning techniques: [gradient descent](#).

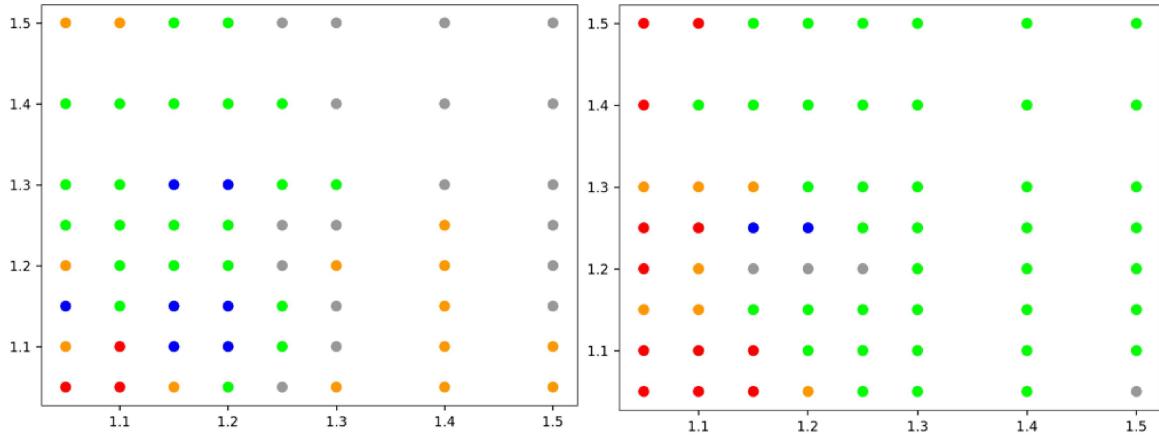


Figure 15.1: Blue/green outperform baseline, orange/red underperform (left: $\omega = 40$, right: $\omega = 60$)

Interesting conclusions:

- There are areas with relatively stable and good performance (green) in the parameter space. There may be many of them, separated by rather large regions of medium or low performance. For gradient descent, as long as it lands in one of them, the results will be good. However, progress within a green zone can be very slow, as the gradient approaches zero. You can stop the descent once deep enough inside a green zone. See how gradient descent works on datasets like this one, without [loss function](#), in section 14.4. An alternative method to find the optimum is [smart grid search](#).
- You want to find a large green region with good or medium performance, rather than a small one with good or great performance. Also, the boundary of these regions can be very unstable, thus the reason to avoid small ones. Note that large green regions may have red zones inside. Do not get too close, these red zones are like black holes. Some of the best parameters (the blue dots) are very close to some of the worst (red dots). Avoid these blue dots to reduce [overfitting](#). This is typical in many problems, where

green zones are **basins of attraction** in the underlying chaotic **dynamical system** acting behind the scenes, surrounded by basins of repulsion. For more details, see my book [9] on this topic.

- The somewhat chaotic boundaries of the various regions can be detected using the **interpolation** algorithm described in section 14.3. This algorithm generalizes to higher dimensions.

Now I focus on the problem addressed in this chapter: designing S&P 500 trading strategies that outperform the baseline (staying long the entire period). The idea consists of exiting (selling) and re-entering (buying) at the right times, including for the first entry, over a long time period, ranging from 3 to 10 years. Depending on the parameters β, σ, ω to be explained later, the number of exits ranges from one or two, to a dozen. At the end of the pre-specified time period, two scenarios are possible: either you have already exited, or you still have an open position. The latter is called *holding*. The right times to exit or re-enter are determined by price variations in the index, your accumulated ROI, the price at your last exit, the time elapsed since your last exit, and the new price. It is well-known that this index, while going up on average, experiences sharp drops – sometimes prolonged – and rebounds that can be leveraged.

15.1 Non-standard strategies

It is difficult to successfully arbitrage the stock market due to the large number of participants competing against you. Staying long on the S&P 500 index is one of the most effective strategies, even outperforming many if not most professional traders, in the long run. In order to do better while minimizing competition, you need to use strategies that Wall Street professionals must avoid. For instance, keeping cash for extended periods of time (sometimes for years in a row) to be able to jump in at the right time with no advance notice – after a massive crash – then buy and sell during a short window following the crash to leverage the resulting volatility, to finally have a stable long position acquired at a steeply discounted price. You then sell the position in question years or months later when its value has massively increased following the slow or fast rebound. Then you repeat the cycle.

The strategies discussed here are inspired by the above principle. Each strategy has a specific parameter vector (β, σ, ω) and its own time frame, starting at t_0 and ending at t_f . The first buy takes place at a time t'_0 in that window, typically with $t'_0 > t_0$, and the last sell at a time t'_f either within that window, or later on if you must still hold your position at time t_f . The annualized compound return is compared to the baseline: buying at t_0 , selling at t_f . If you still hold your position at t_f , the return on your strategy is based on the value of your position at t_f . The starting point to compute your return is t_0 , not t'_0 .

If you are all cash, you buy when the price reaches $\beta \cdot p_{\min}(\omega)$, where $p_{\min}(\omega)$ is the bottom price observed in the last w days preceding the purchase. If holding (active position), you sell when the price hits the $\sigma \cdot p_{\text{buy}}$ threshold, where p_{buy} is the price you paid per unit in your last purchase. Much of the work consists of testing various (β, σ, ω) that works well across various time frames $[t_0, t_f]$, and that are stable. That is, not sensitive to small variations. Now that everything is in place, I make the link to the names used in the Python code in section 15.3.

- In the code, $t_0, t_f, \beta, \sigma, \omega$ are respectively named `init`, `end`, `buy_param`, `sell_param`, `windowLow`, and stored in the `params` dictionary. See lines 121–125 in the code. One extra parameter `rnd` is discussed later. For t'_0 and t'_f , the corresponding notations are `entry_idate` and `exit_idate`. The prices p_{buy} and p_{\min} are respectively denoted as `buy_price` and `bottom[params]`. Finally, `params` is the key to many key-value tables.
- I tested 4 values for t_0 and 3 for t_f . Thus, for each strategy (β, σ, ω) , the summary statistics are based on $4 \times 3 = 12$ observations. An observation is the result of trading the entire time period $[t_0, t_f]$. See lines 47–48 in the code. Note that instead of specifying t_f , I use `duration`, with $t_f = t_0 + \text{duration}$.
- I tested 8 values for β , 8 for σ , and 2 for ω , leading to $8 \times 8 \times 2 = 128$ strategies, each with 12 observations based on the 12 combinations of t_0, t_f . See lines 49–51 in the code. Thus, the output dataset (see [here](#)) has 128 rows: one per strategy, each featuring averages based on 12 measurements.
- Each daily price is a value between the low and the high of the day, specified by the fixed weight `rnd=0.50`. See line 128 in the code. You should try with different values of `rnd`, to rule out strategies too sensitive to little variations.

The code is somewhat complicated because of the pre-processing step to accelerate many computations. The bottleneck is calculating the minimum price in the moving window consisting of the last w days. It is done once in lines 58–61, rather than for each (β, σ, ω) , speeding up the computations by a factor 128.

Other important metrics attached to a strategy are the number `nbuys` of “buy” trades, and among those trades, the `wins`, that is, when buying at a lower price than the previous sell. See lines 81 and 84 in the code. The hold rate is the chance to still be in a holding position at the end of the time period. In the results, all these numbers are averaged per strategy (β, σ, ω) with the strategy average computed over the 12 time series of trading activity (one for each $[t_0, t_f]$). For a strategy, the success rate is defined as the chance to outperform the “buy and hold” baseline.

15.2 Selecting a strategy based on its features

Figure 15.2 shows the summary results for the top 30 strategies, out of 128. The annualized ROI gain over the baseline ranges from 0.82% to -1.71%. Here the baseline is a 4.97% ROI. The average number of “buys” per strategy ranges from barely above 1 to nearly 11. Among those buys, usually fewer than 50% are “wins”, meaning buying at a lower price than the previous sell. Despite this seemingly low performance, the overall success rate (outperforming baseline) per strategy may be over 50% due to the fact that a few of the wins are spectacular, taking place after a massive price drop.

β	σ	ω	buys	wins	delta return	base return	strategy return	success rate	hold rate
1.20	1.10	40	6.25	3.00	0.82%	4.97%	5.79%	67%	83%
1.05	1.15	40	4.17	2.17	0.71%	4.97%	5.68%	42%	67%
1.15	1.25	60	2.75	1.67	0.62%	4.97%	5.59%	58%	83%
1.15	1.10	40	6.00	2.75	0.58%	4.97%	5.55%	50%	83%
1.15	1.30	40	2.50	1.33	0.57%	4.97%	5.53%	75%	83%
1.20	1.25	60	2.75	1.75	0.56%	4.97%	5.53%	58%	83%
1.15	1.15	40	4.42	2.00	0.55%	4.97%	5.52%	67%	83%
1.20	1.15	40	4.42	1.92	0.51%	4.97%	5.48%	67%	92%
1.20	1.30	40	2.58	1.00	0.51%	4.97%	5.48%	58%	100%
1.25	1.25	60	2.83	1.25	0.49%	4.97%	5.46%	67%	92%
1.30	1.10	60	6.08	2.42	0.48%	4.97%	5.45%	50%	83%
1.25	1.15	60	4.33	1.83	0.47%	4.97%	5.44%	58%	83%
1.30	1.30	60	2.33	0.83	0.47%	4.97%	5.44%	50%	75%
1.30	1.15	60	4.33	1.42	0.45%	4.97%	5.42%	50%	92%
1.20	1.25	40	2.83	1.00	0.42%	4.97%	5.39%	42%	100%
1.20	1.10	60	5.75	3.17	0.42%	4.97%	5.39%	58%	50%
1.20	1.30	60	2.33	1.08	0.42%	4.97%	5.39%	67%	75%
1.20	1.05	40	11.08	6.50	0.38%	4.97%	5.35%	42%	83%
1.10	1.30	40	2.42	1.08	0.38%	4.97%	5.35%	58%	83%
1.25	1.30	60	2.33	0.92	0.38%	4.97%	5.35%	50%	75%
1.10	1.15	40	4.17	2.08	0.38%	4.97%	5.35%	25%	75%
1.15	1.25	40	2.67	0.92	0.37%	4.97%	5.34%	33%	83%
1.25	1.05	60	10.58	4.25	0.34%	4.97%	5.31%	42%	50%
1.15	1.50	40	1.92	0.83	0.33%	4.97%	5.30%	50%	100%
1.40	1.25	60	2.75	1.17	0.31%	4.97%	5.28%	58%	92%
1.30	1.25	60	2.75	1.17	0.31%	4.97%	5.28%	58%	92%
1.15	1.40	40	2.17	0.83	0.29%	4.97%	5.26%	67%	92%
1.20	1.40	40	2.17	0.83	0.29%	4.97%	5.26%	67%	92%
1.50	1.25	60	2.75	1.00	0.27%	4.97%	5.24%	58%	92%
1.25	1.30	40	2.50	0.92	0.27%	4.97%	5.24%	42%	100%
1.30	1.05	60	10.83	4.42	0.23%	4.97%	5.20%	50%	67%
1.20	1.50	40	1.92	1.00	0.22%	4.97%	5.19%	50%	100%
1.40	1.40	60	2.17	0.83	0.22%	4.97%	5.19%	58%	92%

Figure 15.2: Average stats for each strategy (β, σ, ω)

Not surprisingly, the chances to finish with an active position (holding) is rather high due to the small volume of trading activity, favoring long over short-term holding positions. See rightmost column in Figure 15.2. It would be interesting to see what proportion of time is spent on holding a position, versus being all cash.

How to choose a strategy, that is, β, σ and ω ? A large number of “buys” with 40% or more that are “wins” is an indication of robustness. Another benefit is that you are more frequently in a cash position. This can be useful if you face some sudden emergencies. However, frequent trades may be subject to short term capital gains. You also want to avoid good strategies (green or blue dot) with parameters β, σ, ω too close to a red dot in Figure 15.1.

Finally, most of the tests were performed on the last 20 years of historical data. Earlier patterns are slightly different and more erratic. Also, you can play with multiple strategies in parallel. The strategies can be used

to identify entry and exit points in the S&P 500 index. It would be interesting the break down performance based on the length $t_f - t_0$ of the time window

15.3 Python code and dataset: 40 years' worth of historical data

The Python code `spx500.py` is on GitHub in my “Statistical Optimization” repository: click [here](#) to access it. The input dataset `spx500.txt` and output results `spx500-results.xlsx` are in the same repository, respectively [here](#) and [here](#). The input data comes from Yahoo Finance: you can also download it from [here](#), and also access data for individual stocks.

```

1 import numpy as np
2 import pandas as pd
3 from datetime import datetime
4 import matplotlib.pyplot as plt
5 import matplotlib as mpl
6
7
8 #--- [1] Read data
9
10 data = pd.read_csv("spx500.txt")
11
12 def date_to_int(date, start_date):
13     xstart_date = datetime.strptime(start_date, '%b %d %Y')
14     xdate = datetime.strptime(date, '%b %d %Y')
15     date_int = str(xdate - xstart_date)
16     date_int = date_int.split(' ')[0]
17     if date_int == '0:00:00':
18         date_int = 0
19     else:
20         date_int = int(date_int)
21     return(date_int)
22
23 arr_date = data['Date']
24 arr_low = data['Low']
25 arr_high = data['High']
26 arr_open = data['Open']
27 arr_idate = []
28 arr_tdate = []
29 nobs = len(arr_date)
30
31 for k in range(nobs):
32     arr_idate.append(date_to_int(arr_date[k], arr_date[0]))
33     tdate = datetime.strptime(arr_date[k], '%b %d %Y')
34     arr_tdate.append(tdate)
35
36 data.insert(1, 'iDate', arr_idate, True)
37
38 print(data.head())
39 print(data.columns.values)
40 print(data.shape)
41 # features: 'Date','iDate','Open','High','Low','Close','Adj_Close','Volume'
42
43 #--- [2] Initializations for main loop
44
45 h_params = {}
46
47 l_init = (4000, 5000, 6000, 7000)
48 l_duration = (1000, 2000, 3000)
49 l_windowLow = (40, 60)
50 l_buy_param = (1.05, 1.1, 1.15, 1.2, 1.25, 1.3, 1.4, 1.5)
51 l_sell_param = (1.05, 1.1, 1.15, 1.2, 1.25, 1.3, 1.4, 1.5)
52
53 h_max = {}
54 h_min = {}
55
56 #- [2.1] preprocessing to speed up computations
57
58 for k in range(nobs):
59     for windowLow in l_windowLow:
60         if k >= windowLow:
61             h_min[(k, windowLow)] = min(arr_low[k-windowLow:k-1])
62
63 #- [2.2] create table of parameter sets (stored as key in h_params)
```

```

64
65 for buy_param in l_buy_param:
66     for sell_param in l_sell_param:
67         for windowLow in l_windowLow:
68             for init in l_init:
69                 for duration in l_duration:
70                     end = init + duration
71                     key = (init,end,buy_param,sell_param,windowLow)
72                     h_params[key] = 1
73
74 #-- [2.3] Create main hash tables
75 #       The index (same in all tables) is the parameter set
76
77 bottom   = {}
78 hold     = {} # true if we have open position
79 value    = {}
80 max_hold = {}
81 nbuys    = {} # number buys done during whole period
82 entry_price = {} # buy price on first buy
83 sell_price = {} # sell price on last sell, before current buy
84 wins     = {} # a win is buying a lower price than last sell
85
86 init_price = {} # price when trading period starts
87 buy_price = {} # buy price
88 end_price = {} # price when trading period ends
89 entry_idate = {} # date of first buy
90 exit_idate = {} # date of last sell if out, otherwise end_idate
91 init_idate = {} # date when trading period starts
92 end_idate = {} # date when trading period ends
93 start    = {}
94
95 for params in h_params:
96
97     bottom[params] = arr_high[0]
98     hold[params]   = False
99     value[params]  = 0
100    max_hold[params] = 0
101    nbuys[params]  = 0
102    entry_price[params] = -1
103    sell_price[params] = -1
104    wins[params]   = 0
105
106
107 #--- [3] Main loop
108
109 for k in range(nobs):
110     # loop over all trading days
111
112     idate = arr_idate[k]
113     price_high = arr_high[k]
114     price_low = arr_low[k]
115     if k % 100 == 0:
116         print(arr_date[k])
117
118     for params in h_params:
119         # update trading stats for each param set in parallel
120
121         init      = params[0]
122         end       = params[1]
123         buy_param = params[2]
124         sell_param = params[3]
125         windowLow = params[4]
126         rnd       = 0.5
127
128         price = rnd * price_high + (1-rnd) * price_low
129
130         if k == init:
131             init_price[params] = price
132             init_idate[params] = idate
133         elif k == end:
134             end_price[params] = price
135             end_idate[params] = idate
136
137         if k >= windowLow and price < h_min[(k, windowLow)]:
138             bottom[params] = price
139

```

```

140     if k >= init and not hold[params] and k < end:
141         if price < buy_param * bottom[params]:
142             # buy
143             buy_price[params] = price
144             if value[params] == 0:
145                 # first purchase
146                 value[params] = price
147                 entry_price[params] = price
148                 entry_idate[params] = idate
149                 if price < init_price[params]:
150                     # first buy at lower price than init price
151                     wins[params] += 1
152             elif buy_price[params] < sell_price[params]:
153                 wins[params] += 1
154             start[params] = idate
155             hold[params] = True
156             nbuys[params] += 1
157
158     elif hold[params] and k < end:
159         span = idate-start[params]
160         if span > max_hold[params]:
161             max_hold[params] = span
162         if price > sell_param * buy_price[params]:
163             # sell
164             sell_price[params] = price
165             value[params] *= price/buy_price[params]
166             hold[params]= False
167             exit_idate[params] = idate
168
169 #---- [4] Summary stats for each param set
170
171 # group params in h_params by (init, end)
172 # stored average stats by grouped params, in arr_local
173
174 hash_performance = {}
175 hash_count = {}
176
177 for params in h_params:
178     if hold[params]:
179         value[params] *= end_price[params]/buy_price[params]
180         exit_idate[params] = end_idate[params]
181     else:
182         if end_price[params] < sell_price[params]:
183             wins[params] += 1
184     duration1 = end_idate[params] - init_idate[params] + 1
185     duration2 = exit_idate[params] - entry_idate[params] + 1
186     R_market = end_price[params] / init_price[params]
187     R_strategy = value[params] / entry_price[params]
188     adj_R_market = 100*(R_market**((365/duration1) - 1))
189     adj_R_strategy = 100*(R_strategy**((365/duration1) - 1))
190     ratio = R_strategy / R_market # reinvest all to compound return
191     performance = adj_R_strategy - adj_R_market
192     if performance > 0:
193         success = 1
194     else:
195         success = 0
196
197     key = (params[2], params[3], params[4])
198     arr_local = [nbuys[params], wins[params],
199                  performance, adj_R_market, adj_R_strategy,
200                  duration1, duration2, success, hold[params]]
201
202     if key in hash_performance:
203         arr_local2 = hash_performance[key]
204         for idx in range(len(arr_local2)):
205             arr_local2[idx] += arr_local[idx]
206         hash_performance[key] = arr_local2
207         hash_count[key] += 1
208     else:
209         hash_performance[key] = arr_local
210         hash_count[key] = 1
211
212 #- [4.1] to get averages for each param set, divide sums by sample size cnt
213
214 cnt = hash_count[key]
```

```

216
217     for key in hash_performance:
218         arr_avg = hash_performance[key]
219         for idx in range(len(arr_avg)):
220             arr_avg[idx] /= cnt
221
222
223     #--- [5] Print results
224
225     OUT = open("spx500-results2.txt", "w")
226
227     labels="beta\tsigma\tomega\tsample size\tbuys\tbuys below last sell\tDelta return"
228     labels+= "base return\ttrading return\tperiod (days)\tactive peiod\tsuccess rate\thold rate"
229     OUT.write(labels + "\n")
230
231     for key in hash_count:
232         strout = ""
233         for param in key:
234             strout += str(param) + "\t"
235         cnt = hash_count[key]
236         strout += str(cnt) + "\t"
237         arr_avg = hash_performance[key]
238         for idx in range(len(arr_avg)):
239             strout += str(arr_avg[idx]) + "\t"
240         strout += "\n"
241     OUT.write(strout)
242
243     OUT.close()
244
245
246     #--- [6] Visualizations
247
248     hash_xy = {}
249     hash_color = {}
250     for value in l_windowLow:
251         hash_xy[value] = []
252         hash_color[value] = []
253
254     for key in hash_count:
255         cnt = hash_count[key]
256         arr_avg = hash_performance[key]
257         for params in key:
258             buy_param = key[0]
259             sell_param = key[1]
260             window_low = key[2]
261             performance = arr_avg[2] # performance
262             if performance > 0.5:
263                 color = [0, 0, 1] # blue
264             elif performance > 0.1:
265                 color = [0, 1, 0] # lightgreen
266             elif performance > -0.1:
267                 color = [0.6, 0.6, 0.6] # lightgray
268             elif performance > -0.5:
269                 color = [1, 0.6, 0] # orange
270             else:
271                 color = [1, 0, 0] # red
272             local_arr = hash_xy[window_low]
273             local_arr.append([buy_param, sell_param])
274             local_arr = hash_color[window_low]
275             local_arr.append(color)
276
277     mpl.rcParams['axes.linewidth'] = 0.5
278     plt.rcParams['xtick.labelsize'] = 9
279     plt.rcParams['ytick.labelsize'] = 9
280
281     for window_low in hash_xy:
282         z = np.array(hash_xy[window_low])
283         z = np.transpose(z)
284         color = hash_color[window_low]
285         plt.scatter(z[0], z[1], c = color)
286         plt.show()

```
